

Invert goals: in-context learning vs weight updates

Grey bars score 0 % equivalence: the model produces syntactically valid PDDL with the wrong goal in every case.

Equivalence (%)

70
60
50
40
30
20
10
0

0.0%

base
(0-shot)

0.0%

4-shot
type-matched
(n=50)

59.6%

fine-tuned
(in-domain)

